Statistics

Collection of Formulas

Contents

1 Descriptive Statistics

1.1 Summary Statistics

1.1.1 Location

Mode   Most frequent value of x_i. Two or more modes are possible (bimodal).

Median
$$\tilde{x}_{0.5} = \begin{cases} x_{((n+1)/2)} & \text{falls } n \text{ ungerade} \\ \frac{1}{2}(x_{(n/2)} + x_{(n/2+1)}) & \text{falls } n \text{ gerade} \end{cases}$$

Quantile
$$\tilde{x}_{\alpha} = \begin{cases} x_{(k)} & \text{falls } n\alpha \notin \mathbb{N} \\ \frac{1}{2}(x_{(n\alpha)} + x_{(n\alpha+1)}) & \text{falls } n\alpha \text{ ganzzahlig} \end{cases}$$
with
$$k = \min x \in \mathbb{N}, \quad x > n\alpha$$

Minimum/Maximum
$$x_{\min} = \min_{i \in \{1,\dots,N\}}(x_i) \qquad x_{\max} = \max_{i \in \{1,\dots,N\}}(x_i)$$

Arithmetic Mean
$$\bar{x} = \frac{1}{n}\sum_{i=1}^{n} x_i$$

Estimates the expectation $\mu = E[X]$ (first moment).
Calculation Rules:
  ⋆ $E(a + b \cdot X) = a + b \cdot E(X)$
  ⋆ $E(X \pm Y) = E(X) \pm E(Y)$

Geometric Mean
$$\bar{x}_G = \sqrt[n]{\sum_{i=1}^{n} x_i}$$

For growth factors: $\bar{x}_G = \sqrt[n]{\frac{B_n}{B_0}}$

Harmonic Mean
$$\bar{x}_H = \frac{\sum_{i=1}^{n} w_i}{\sum_{i=1}^{n} \frac{w_i}{x_i}}$$

1.1.2 Dispersion

Range
$$R = x_{(n)} - x_{(1)}$$

Interquartile Range
$$d_Q = \tilde{x}_{0.75} - \tilde{x}_{0.25}$$

(Empirical) Variance
$$s^2 = \frac{1}{n}\sum_{i=1}^{n}(x_i - \bar{x})^2 = \frac{1}{n}\sum_{i=1}^{n} x_i^2 - \bar{x}^2$$

Estimates the second centralized moment.
Calculation Rules:
  ⋆ $Var(aX + b) = a^2 \cdot Var(X)$

  ⋆ $Var(X \pm Y) = Var(X) + Var(Y) + 2Cov(X, Y)$

(Empirical) Standard Deviation
$$s = \sqrt{s^2}$$

Coefficient of Variation
$$\nu = \frac{s}{\bar{x}}$$

Average Absolute Deviation
$$e = \frac{1}{n}\sum_{i=1}^{n} |x_i - \bar{x}|$$

Estimates the first absolute centralized moment.

1.1.3 Concentration

Gini Coefficient
$$G = \frac{2\sum_{i=1}^{n} ix_{(i)} - (n+1)\sum_{i=1}^{n} x_{(i)}}{n\sum_{i=1}^{n} x_{(i)}} = 1 - \frac{1}{n}\sum_{i=1}^{n}(v_{i-1} + v_i)$$
with
$$u_i = \frac{i}{n}, \quad v_i = \frac{\sum_{j=1}^{i} x_{(j)}}{\sum_{j=1}^{i} x_{(j)}} \qquad (u_0 = 0, \quad v_0 = 0)$$

These are also the values for the Lorenz curve.

Range: $0 \leq G \leq \frac{n-1}{n}$

Lorenz-Münzner Coefficient (normed G)
$$G^+ = \frac{n}{n-1}G$$

Range: $0 \leq G^+ \leq 1$

1.1.4 Shape

(Empirical) Skewness

$$\nu = \frac{n}{(n-1)(n-2)} \sum_{i=1}^{n} \left(\frac{x_i - \bar{x}}{s} \right)^3$$

Estimates the third centralized moment, scaled with $(\sigma^2)^{\frac{2}{3}}$

(Empirical) Kurtosis

$$k = \left[n(n+1) \cdot \sum_{i=1}^{n} \left(\frac{x_i - \bar{x}}{s} \right)^4 - 3(n-1) \right] \cdot \frac{n-1}{(n-2)(n-3)} + 3$$

Estimates the fourth centralized moment, scaled with $(\sigma^2)^2$

Excess

$$\gamma = k - 3$$

1.1.5 Dependence

for two nominal variables

χ^2-Statistic

$$\chi^2 = \sum_{i=1}^{k} \sum_{j=1}^{l} \frac{(n_{ij} - \frac{n_{i+}n_{+j}}{n})^2}{\frac{n_{i+}n_{+j}}{n}} = n \left(\sum_{i=1}^{k} \sum_{j=1}^{l} \frac{n_{ij}^2}{n_{i+}n_{+j}} - 1 \right)$$

Range: $0 \leq \chi^2 \leq n(\min(k,l) - 1)$

Phi-Coefficient

$$\Phi = \sqrt{\frac{\chi^2}{n}}$$

Range: $0 \leq \Phi \leq \sqrt{\min(k,l) - 1}$

Cramér's V

$$V = \sqrt{\frac{\chi^2}{\min(k,l) - 1}}$$

Range: $0 \leq V \leq 1$

Contingency Coefficient C

$$C = \sqrt{\frac{\chi^2}{\chi^2 + n}}$$

Range: $0 \leq C \leq \sqrt{\frac{\min(k,l) - 1}{\min(k,l)}}$

Corrected Contingency Coefficient C_{corr}

$$C_{corr} = \sqrt{\frac{\min(k,l)}{\min(k,l) - 1}} \cdot \sqrt{\frac{\chi^2}{\chi^2 + n}}$$

Range $0 \leq C_{corr} \leq 1$

Odds-Ratio

$$OR = \frac{ad}{bc} = \frac{n_{ii}n_{jj}}{n_{ij}n_{ji}}$$

Range: $0 \leq OR < \infty$

for two ordinal variables

Gamma (Goodman and Kruskal)

$$\gamma = \frac{K - D}{K + D}$$

$K = \sum_{i<m} \sum_{j<n} n_{ij}n_{mn}$ Number of concordant pairs
$D = \sum_{i<m} \sum_{j>n} n_{ij}n_{mn}$ Number of reversed pairs

Range: $-1 \leq \gamma \leq 1$

Kendall's τ_b

$$\tau_b = \frac{K - D}{\sqrt{(K + D + T_X)(K + D + T_Y)}}$$

with
$T_X = \sum_{i=m} \sum_{j<n} n_{ij}n_{mn}$ Number of ties w.r.t. X
$T_Y = \sum_{i<m} \sum_{j=n} n_{ij}n_{mn}$ Number of ties w.r.t. Y

Range: $-1 \leq \tau_b \leq 1$

Kendall's/Stuart's τ_c

$$\tau_c = \frac{2 \min(k,l)(K - D)}{n^2(\min(k,l) - 1)}$$

Range: $-1 \leq \tau_c \leq 1$

Spearman's Rank Correlation Coefficient

$$\rho = \frac{n(n^2 - 1) - \frac{1}{2}\sum_{j=1}^{J} b_j(b_j^2 - 1) - \frac{1}{2}\sum_{k=1}^{K} c_k(c_k^2 - 1) - 6\sum_{i=1}^{n} d_i^2}{\sqrt{n(n^2 - 1) - \sum_{j=1}^{J} b_j(b_j^2 - 1)} \sqrt{n(n^2 - 1) - \sum_{k=1}^{K} c_k(c_k^2 - 1)}}$$

or

$$\rho = \frac{s_{rg_x rg_y}}{\sqrt{s_{rg_x rg_x} s_{rg_y rg_y}}}$$

Without ties:

$$\rho = 1 - \frac{6\sum_{i=1}^{n} d_i^2}{n(n^2 - 1)}$$

with
$d_i = R(x_i) - R(y_i)$ rank difference

Range: $-1 \leq \rho \leq 1$

for two metric variables

Correlation Coefficient (Bravais-Pearson)

$$r = \frac{S_{xy}}{\sqrt{S_{xx}S_{yy}}} = \frac{s_{xy}}{\sqrt{s_{xx}s_{yy}}}$$

with
$S_{xy} = \sum_{i=1}^{n} (x_i - \bar{x})^2(y_i - \bar{y})^2$ or $s_{xy} = \frac{S_{xy}}{n}$
$S_{xx} = \sum_{i=1}^{n} (x_i - \bar{x})^2$ or $s_{xx} = \frac{S_{xx}}{n}$
$S_{yy} = \sum_{i=1}^{n} (y_i - \bar{y})^2$ or $s_{yy} = \frac{S_{yy}}{n}$

Range: $-1 \leq r \leq 1$

1.2   Tables

1.3   Diagrams

1.3.1   Histogram
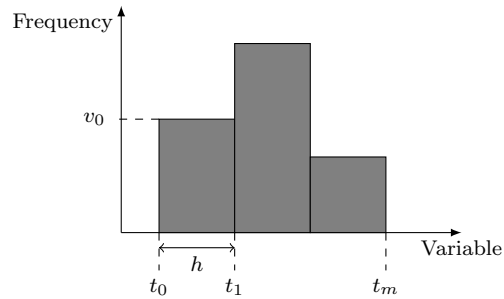
sample: $X = \{x_1, x_2, ...; x_n\}$

k-th bin: $B_k = [t_k, t_{k+1}), k = \{0, 1, ..., m-1\}$

Number of observations in the k-th bin: v_k

bin width: $h = t_{k+1} - t_k, \forall k$

Scott's Rule

$$h^* \approx 3.5\sigma n^{-\frac{1}{3}}$$

For approximately normal distributed data (min. MSE)

2 Probability

2.1 Combinatorics

		without replacement	with replacement
Permutations		$n!$	$\frac{n!}{n_1! \cdots n_s!}$
Combinations:	without order	$\binom{n}{m}$	$\binom{n+m-1}{m}$
	with order	$\binom{n}{m} m!$	n^m

with:
$$n! = n \cdot (n-1) \cdot \ldots \cdot 1$$
$$\binom{n}{m} = \frac{n!}{m!(n-m)!}$$

2.2 Probability Theory

Laplace
$$P(A) = \frac{|A|}{|\Omega|}$$

Kolmogorov Axioms   mathematical definition of probability

(1) $0 \leq P(A) \leq 1 \quad \forall A \in \mathcal{A}$
(2) $P(\Omega) = 1$
(3) $P(\bigcup_{i=1}^{\infty} A_i) = \sum_{i=1}^{\infty} P(A_i)$
      $\forall A_i \in \mathcal{A}, i = 1, ..., \infty$ with $A_i \cap A_j = \emptyset$ for $i \neq j$

Implications:

- $P(\bar{A}) = 1 - P(A)$

- $P(\emptyset) = 0$

- $P(A \cup B) = P(A) + P(B) - P(A \cap B)$

- $A \subseteq B \Rightarrow P(A) \leq P(B)$

- $P(B) = \sum_{i=1}^{n} P(B \cap A_i)$, for $A_i, ..., A_n$ complete decomposition of Ω into pairwise disjoint events

Probability (Mises)   frequentist definition of probability
$$P(A) = \lim_{n \to \infty} \frac{n_A(n))}{n}$$
with n repetitions of a random experiment and $n_A(n)$ events A

Conditional Probability
$$P(A|B) = \frac{P(A \cap B)}{P(B)} \qquad \text{für } P(B) > 0$$
$$\Rightarrow P(A \cap B) = P(B|A)P(A) = P(A|B)P(B)$$

Law of Total Probability
$$P(B) = \sum_{i=1}^{n} P(B|A_i)P(A_i)$$

Bayes' Theorem
$$P(A|B) = \frac{P(B|A)P(A)}{P(B)} \qquad \text{for } P(A), P(B) > 0$$

Stochastic Independence

A, B independent $\Leftrightarrow P(A \cap B) = P(A) + P(B)$

X, Y independent $\Leftrightarrow f_{XY}(x,y) = f_X(x) \cdot f_Y(y) \qquad \forall x, y$

2.3 Random Variables/Vectors

Random Variables $\in \mathbb{R}$

Definition
$$Y : \Omega \to \mathbb{R}$$
The Subset of possible values for $\mathbb{R}$ is called support.
Notation: Realisations of Y are depicted with lower case letters.
$Y = y$ means, that y is the realisation of Y.

Discrete and Continuous Random Variables
If the support is uncountably infinite, the random variable is called *continuous*, otherwise it is called *discrete*.

- **Density $f(\cdot)$:**

  For continuous variables: $P(Y \in [a,b]) = \int_a^b f_Y(y)dy$

  For discrete variables the density (and other functions) can be depicted like the corresponding function for continuous variables, if the notation is extended as follows:
  $\int_{-\infty}^{y} f_Y(\tilde{y})d\tilde{y} := \sum_{k: k \leq y} P(Y = k)$. This notation is used.

- **Cumulative Distribution Function $F(\cdot)$:**
  $F_Y(y) = P(Y \leq y)$

  Relationship:
  $$F_Y(y) = \int_{-\infty}^{y} f_Y(\tilde{y})d\tilde{y}$$

Moments

- **Expectation (1. Moment):** $\mu = E(Y) = \int y f_Y(y)dy$

- **Variance (2. centralized Moment):**
  $\sigma^2 = Var(Y) = E(\{Y - E(Y)\}^2) = \int (y - E(Y))^2 f(y)dy$
  Note: $E(\{Y - \mu\}^2) = E(Y^2) - \mu^2$

  > Proof:
  > $E(\{Y-\mu\}^2) = E(Y^2 - 2Y\mu + \mu^2) = E(Y^2) - 2\mu^2 + \mu^2 = E(Y^2) - \mu^2$

- **kth Moment**: $E(Y^k) = \int y^k f_Y(y)dy$,
  k. centralized Moment: $E(\{Y - E(Y)\}^k)$

Moment Generating Function

$$M_Y(t) = \mathrm{E}(e^{tY})$$

with $\left.\frac{\partial^k M_Y(t)}{\partial t^k}\right|_{t=0} = \mathrm{E}(Y^k)$

Cumulant Generating Function $K_Y(t) = \log M_Y(t)$

A random variable is uniquely defined by its moment generating function and vice versa (as long as moments and cumulants are finite).

Random Vectors $\in \mathbb{R}^q$

Density and Cumulative Distribution Function

$$F(y_1, ..., y_q) = P(Y_1 \le y_1, ..., Y_q \le y_q)$$

$$P(a_1 \le Y_1 \le b_1, ..., a_q \le Y_q \le b_q)$$
$$= \int_{a_1}^{b_1} ... \int_{a_q}^{b_q} f(y_1, .., y_q) dy_1 ... dy_q$$

Marginal Density

$$f_{Y_1}(y_1) = \int_{-\infty}^{\infty} ... \int_{-\infty}^{\infty} f(y_1, ..., y_k) dy_2 ... dy_k$$

Conditional Density

$$f_{Y_1|Y_2}(y_1|y_2) = \frac{f(y_1, ..., y_2)}{f(y_2)} \text{ for } f(y_2) > 0$$

Iterated Expectation

$$\mathrm{E}(Y) = \mathrm{E}_X(\mathrm{E}(Y|X))$$

Proof:

$$\mathrm{E}(Y) = \int y f(y) dy = \int \int y f(y|x) dy f_X(x) dx = \mathrm{E}_X(\mathrm{E}(Y|X))$$

$$\mathrm{Var}(Y) = \mathrm{E}_X(\mathrm{Var}(Y|X)) + \mathrm{Var}_X(\mathrm{E}(Y|X))$$

Proof:

$$\mathrm{Var}(Y) = \int (y - \mu_Y)^2 f(y) dy$$
$$= \int (y - \mu_Y)^2 f(y|x) f(x) dy dx$$
$$= \int (y - \mu_{Y|x} + \mu_{Y|x} - \mu_Y)^2 f(y|x) f(x) dy dx$$
$$= \int (y - \mu_{Y|x})^2 f(y|x) f(x) dy dx +$$
$$\int (\mu_{Y|x} - \mu_Y)^2 f(y|x) f(x) dy dx +$$
$$2 \int (y - \mu_{Y|x})(\mu_{Y|x} - \mu_Y) f(y|x) f(x) dy dx$$
$$= \int \mathrm{Var}(Y|x) f(x) dx + \int (\mu_{Y|x} - \mu_Y)^2 f(x) dx$$
$$= \mathrm{E}_X(\mathrm{Var}(Y|X)) + \mathrm{Var}_X(\mathrm{E}(Y|X))$$

2.4 Probability Distributions

2.4.1 Discrete Distributions

Discrete Uniform

$$Y \sim \mathrm{U}(\{y_1, ..., y_k\}), \ y \in \{y_1, ..., y_k\}$$
$$P(Y = y_i) = \frac{1}{k}, \ i = 1, ..., k$$
$$\mathrm{E}(Y) = \frac{k+1}{2}, \ \mathrm{Var}(Y) = \frac{k^2 - 1}{12}$$

Binomial   Successes in independent trials

$$Y \sim \mathrm{Bin}(n, \pi) \text{ with } n \in \mathbb{N}, \pi \in [0, 1], \ y \in \{0, ..., n\}$$
$$P(Y = y|\lambda) = \binom{n}{y} \pi^k (1 - \pi)^{n-y}$$
$$\mathrm{E}(Y|\pi, n) = n\pi, \ \mathrm{Var}(Y|\pi, n) = n\pi(1 - \pi)$$

Poisson   Counting model for rare events
only one event at a time, no autocorrelation, mean number of events over time is constant and proportional to length of the considered time interval

$$Y \sim \mathrm{Po}(\lambda) \text{ with } \lambda \in [0, +\infty], \ y \in \mathbb{N}_0$$

$$P(Y = y|\lambda) = \frac{\lambda^y exp^{-\lambda}}{y!}$$
$$\mathrm{E}(Y|p) = \lambda, \ \mathrm{Var}(Y|p) = \lambda$$

The model tends to overestimate the variance (Overdispersion).
Approximation of the Binomial for small p

Geometric

$$Y \sim \mathrm{Geom}(\pi) \text{ with } \pi \in [0, 1], \ y \in \mathbb{N}_0$$
$$P(Y = y|\pi) = \pi(1 - \pi)^{y-1}$$
$$\mathrm{E}(Y|\pi) = \frac{1}{\pi}, \ \mathrm{Var}(Y|\pi) = \frac{1 - \pi}{\pi^2}$$

Negative Binomial

$$Y \sim \mathrm{NegBin}(\alpha, \beta) \text{ with } \alpha, \beta \ge 0, \ y \in \mathbb{N}_0$$
$$P(Y = y|\alpha, \beta) = \binom{\alpha + y - 1}{\alpha - 1} \left(\frac{\beta}{\beta - 1}\right)^{\alpha} \left(\frac{1}{\beta + 1}\right)^y$$
$$\mathrm{E}(Y|\alpha, \beta) = \frac{\alpha}{\beta}, \ \mathrm{Var}(Y|\alpha, \beta) = \frac{\alpha}{\beta^2}(\beta + 1)$$

2.4.2 Continuous Distributions

Continuous Uniform

$$Y \sim \mathrm{U}(a,b) \text{ with } \alpha, \beta \in \mathbb{R}, a \le b, \ y \in [a,b]$$

$$p(y|a,b) = \frac{1}{b-a}$$

$$\mathrm{E}(Y|a,b) = \frac{a+b}{2}, \ \mathrm{Var}(Y|a,b) = \frac{(b-a)^2}{12}$$

Univariate Normal   symmetric with μ and σ^2

$$Y \sim \mathrm{N}(\mu,\sigma^2) \text{ with } \mu \in \mathbb{R}, \sigma^2 > 0, \ y \in \mathbb{R}$$

$$p(y|\mu,\sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y-\mu)^2}{2\sigma^2}\right)$$

$$\mathrm{E}(Y|\mu,\sigma^2) = \mu, \ \mathrm{Var}(Y|\mu,\sigma^2) = \sigma^2$$

Multivariate Normal   symmetric with μ_i and Σ

$$Y \sim \mathrm{N}(\mu,\Sigma) \text{ with } \mu \in \mathbb{R}^d, \Sigma \in \mathbb{R}^{d \times d} s.p.d., \ y \in \mathbb{R}^d$$

$$p(y|\mu,\Sigma) = (2\pi)^{-\frac{d}{2}} \det(\Sigma)^{-\frac{1}{2}} \exp\left(-\frac{1}{2}(y-\mu)^T \Sigma^{-1}(y-\mu)\right)$$

$$\mathrm{E}(Y|\mu,\Sigma) = \mu, \ \mathrm{Var}(Y|\mu,\Sigma) = \Sigma$$

Log-Normal

$$Y \sim \mathrm{LogN}(\mu,\sigma^2) \text{ eith } \mu \in \mathbb{R}, \sigma^2 > 0, \ y > 0$$

$$p(y|\mu,\sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}y} \exp\left(-\frac{(\log y - \mu)^2}{2\sigma^2}\right)$$

$$\mathrm{E}(Y|\mu,\sigma^2) = \exp(\mu + \frac{\sigma^2}{2}),$$

$$\mathrm{Var}(Y|\mu,\sigma^2) = \exp(2\mu + \sigma^2)(\exp(\sigma^2)-1)$$

Relationship: $\log(Y) \sim \mathrm{N}(\mu,\sigma^2) \Rightarrow Y \sim \mathrm{LogN}(\mu,\sigma^2)$

non-standardized Student's t   statistical Tests for μ with unknown (estimated) variance and ν degrees of freedom

$$Y \sim \mathrm{t}_\nu(\mu,\sigma^2) \text{ with } \mu \in \mathbb{R}, \sigma^2, \nu > 0, \ y \in \mathbb{R}$$

$$p(y|\mu,\sigma^2,\nu) = \frac{\Gamma\left(\frac{\nu+1}{2}\right)}{\Gamma(\frac{\nu}{2})\Gamma(\sqrt{\nu\pi}\sigma)}\left(1 + \frac{(y-\mu)^2}{\nu\sigma^2}\right)^{-\frac{\nu+1}{2}}$$

$$\mathrm{E}(Y|\mu,\sigma^2,\nu) = \mu \text{ for } \nu > 1,$$

$$\mathrm{Var}(Y|\mu,\sigma^2,\nu) = \sigma^2 \frac{\nu}{\nu-2} \text{ for } \nu > 2$$

Relationship: $Y|\theta \sim \mathrm{N}(\mu,\frac{\sigma^2}{\theta}), \ \theta \sim \mathrm{Ga}(\frac{\nu}{2},\frac{\nu}{2}) \Rightarrow Y \sim \mathrm{t}_\nu(\mu,\sigma)$
$t_\nu(\mu,\sigma^2)$ has heavier tails then the normal distribution.
$t_\infty(\mu,\sigma^2)$ approaches $\mathrm{N}(\mu,\sigma^2)$.

Beta

$$Y \sim \mathrm{Be}(a,b) \text{ with } a,b > 0, \ y \in [0,1]$$

$$p(y|a,b) = \frac{\Gamma(a+b)}{\Gamma(a)\Gamma(b)} y^{a-1}(1-y)^{b-1}$$

$$\mathrm{E}(Y|a,b) = \frac{a}{a+b},$$

$$\mathrm{Var}(Y|a,b) = \frac{ab}{(a+b)^2(a+b+1)},$$

$$\mathrm{mod}(Y|a,b) = \frac{a-1}{a+b-2} \text{ for } a,b > 1$$

Gamma

$$Y \sim \mathrm{Ga}(a,b) \text{ with } a,b > 0, \ y > 0$$

$$p(y|a,b) = \frac{b^a}{\Gamma(a)} y^{a-1} \exp(-by)$$

$$\mathrm{E}(Y|a,b) = \frac{a}{b},$$

$$\mathrm{Var}(Y|a,b) = \frac{a}{b^a},$$

$$\mathrm{mod}(Y|a,b) = \frac{a-1}{b} \text{ for } a \ge 1$$

Inverse-Gamma

$$Y \sim \mathrm{IG}(a,b) \text{ with } a,b > 0, \ y > 0$$

$$p(y|a,b) = \frac{b^a}{\Gamma(a)} y^{-a-1} \exp(-\frac{b}{y})$$

$$\mathrm{E}(Y|a,b) = \frac{b}{a-1} \text{ for } a > 1,$$

$$\mathrm{Var}(Y|a,b) = \frac{b^2}{(a-1)^2(a-2)} \text{ for } a \ge 2,$$

$$\mathrm{mod}(Y|a,b) = \frac{b}{a+1}$$

Relationship: $Y^{-1} \sim \mathrm{Ga}(a,b) \Leftrightarrow Y \sim \mathrm{IG}(a,b)$

Exponential   Time between Poisson events

$$Y \sim \mathrm{Exp}(\lambda) \text{ with } \lambda > 0, \ y \ge 0$$

$$p(y|\lambda) = \lambda \exp(-\lambda y)$$

$$\mathrm{E}(Y|\lambda) = \frac{1}{\lambda}, \ \mathrm{Var}(Y|\lambda) = \frac{1}{\lambda^2}$$

Chi-Squared   squared standard normal random variables with ν degrees of freedom

$$Y \sim \chi^2(\nu) \text{ with } \nu > 0,, \ y \in \mathbb{R}$$

$$p(y|\nu) = \frac{y^{\frac{\nu}{2}-1} e^{-\frac{y}{2}}}{2^{\frac{\nu}{2}} \Gamma\left(\frac{\nu}{2}\right)}$$

$$\mathrm{E}(Y|\nu) = \nu, \ \mathrm{Var}(Y|\nu) = 2\nu$$

2.4.3 Exponential Family

Definition

The exponential family comprises all distributions, whose density can be written as follows:

$$f_Y(y,\theta) = \exp^{t^T(y)\theta - \kappa(\theta)} h(y)$$

with $h(y) \ge 0$, $t(y)$ vector of the canonical statistic, θ as parameter and $\kappa(\theta)$ the normalising constant.

Normalising Constant

$$1 = \int \exp^{t^T(y)\theta} h(y) dy \exp^{-\kappa(\theta)}$$

$$\Leftrightarrow \kappa(\theta) = \log \int \exp^{t^T(y)\theta} h(y)dy$$

$\kappa(\theta)$ is the cumulant generating function, therefore $\frac{\partial \kappa(\theta)}{\partial \theta} = \mathrm{E}(t(Y))$ and $\frac{\partial^2 \kappa(\theta)}{\partial \theta^2} = \mathrm{Var}(t(Y))$

Members

- **Poisson**

- **Geometric**

- **Exponential**

- **Normal** $t(y) = \left(-\frac{y^2}{2}, y\right)^T$, $\theta = \left(\frac{1}{\sigma^2}, \frac{\mu}{\sigma^2}\right)^T$, $h(y) = \frac{1}{\sqrt{2\pi}}$, $\kappa(\theta) = \frac{1}{2}\left(-\log\frac{1}{\sigma^2} + \frac{\mu^2}{\sigma^2}\right)$

- **Gamma**

- **Chi-Squared**

- **Beta**

2.5 Limit Theorems

Law of Large Numbers

Central Limit Theorem

$$Z_n \xrightarrow{d} \mathrm{N}(0, \sigma^2)$$

with $Z_n = \sum_{i=1}^n \frac{Y_i}{\sqrt{n}}$ and Y_i i.i.d. with expectation 0 and variance σ^2

Proof:

For normal random variables $Z \sim \mathrm{N}(\mu, \sigma^2)$: $K_Z(t) = \mu t + \frac{1}{2}\sigma^2 t^2$. The first two derivatives $\left.\frac{\partial^k K_Z(t)}{\partial t^k}\right|_{t=0}$ are μ and σ. All other moments are zero.

For $Z_n = (Y_1 + Y_2 + ... + Y_n)/\sqrt{n}$:

$$\begin{aligned}
M_{Z_n}(t) &= \mathrm{E}\left(e^{t(Y_1 + Y_2 + ... + Y_n)/\sqrt{n}}\right)\\
&= \mathrm{E}\left(e^{tY_1/\sqrt{n}} \cdot e^{tY_2/\sqrt{n}} \cdot ... \cdot e^{tY_n/\sqrt{n}}\right)\\
&= \mathrm{E}\left(e^{tY_1/\sqrt{n}}\right)\mathrm{E}\left(e^{tY_2/\sqrt{n}}\right)...\mathrm{E}\left(e^{tY_n/\sqrt{n}}\right)\\
&= M_Y^n(t/\sqrt{n})
\end{aligned}$$

Analoguously: $K_{Z_n}(t) = nK_Y(t/\sqrt{n})$.

$$\left.\frac{\partial K_{Z_n}(t)}{\partial t}\right|_{t=0} = \frac{n}{\sqrt{n}}\left.\frac{\partial K_Y(t)}{\partial t}\right|_{t=0} = \sqrt{n}\mu$$

$$\left.\frac{\partial^2 K_{Z_n}(t)}{\partial t^2}\right|_{t=0} = \frac{n}{n}\left.\frac{\partial^2 K_Y(t)}{\partial t^2}\right|_{t=0} = \sigma^2$$

Using the Taylor Expansion, we can write $K_{Z_n}(t) = 0 + \sqrt{n}\mu t + \frac{1}{2}\sigma^2 t^2 + ...$, where the terms in ... are tending towards 0 as $n \to \infty$.

Therefore: $K_{Z_n}(t) \xrightarrow{n\to\infty} K_Z(t)$ with $Z \sim \mathrm{N}(\sqrt{n}\mu, \sigma^2)$.

3   Inference

3.1   Method of Moments

The theoretical moments are estimated by their empirical counterparts:

$$\mathrm{E}_{\hat{\theta}_{MM}}(Y^k) = m_k(y_1, ..., y_n)$$

For the exponential family: $\hat{\theta}_{MM} = \hat{\theta}_{ML}$

3.2   Loss Functions

Loss

$$\mathcal{L} : \mathcal{T} \times \Theta \to \mathbb{R}^+$$

with parameter space $\Theta \subset \mathbb{R}$, $t \in \mathcal{T}$ with $t : \mathbb{R}^n \to \mathbb{R}$ a statistic, that estimates the parameter θ, $\mathcal{L}(\theta, \theta) = 0$ holds

- **absolute loss (L1)**: $\mathcal{L}(t, \theta) = |t - \theta|$

- **quadratic loss (L2)**: $\mathcal{L}(t, \theta) = (t - \theta)^2$

As θ is unknown, the loss is a theoretical measure. Additionally, it is the realisation of a random variable as it is dependent on a concrete sample.

Risiko

$$R(t(.), \theta) = \mathrm{E}_\theta \left(\mathcal{L}(t(Y_1, ..., Y_n), \theta) \right)$$
$$= \int_{-\infty}^{\infty} \mathcal{L}(t(Y_1, ..., Y_n), \theta) \prod_{i=1}^{n} f(y_i; \theta) dy_i$$

Minimax Approach

The risk still depends ton the true parameter θ. Tentative estimation: Choose θ, so that the risk is maximal and then $t(.)$, so that the risk is minimized (minimizing the worst case):

$$\hat{\theta}_{minimax} = \arg \min_{t(.)} \left(\max_{\theta \in \Theta} R(t(.); \theta) \right)$$

Mean Squared Error (MSE)

$$MSE(t(.), \theta) = \mathrm{E}_\theta \left(\{ t(Y) - \theta \}^2 \right)$$
$$= \mathrm{Var}_\theta \left(t(Y_1, ..., Y_n) \right) + Bias^2((t(.); \theta)$$

with $Bias(t(.); \theta) = \mathrm{E}_\theta \left(t(Y_1, ..., Y_n) \right) - \theta$

> Proof:
> Let $\mathcal{L}(t, \theta) = (t - \theta)^2$
> $R(t(.), \theta) = \mathrm{E}_\theta (\{ t(Y) - \theta \}^2)$
> $\qquad = \mathrm{E}_\theta (\{ t(Y) - \mathrm{E}_\theta(t(Y)) + \mathrm{E}_\theta(t(Y)) - \theta \}^2)$
> $\qquad = \mathrm{E}_\theta (\{ t(Y) - \mathrm{E}_\theta(t(Y)) \}^2) + \mathrm{E}_\theta (\{ \mathrm{E}_\theta(t(Y)) - \theta \}^2)$
> $\qquad \quad + 2\mathrm{E}_\theta (\{ t(Y) - \mathrm{E}_\theta(t(Y)) \} \{ \mathrm{E}_\theta(t(Y)) - \theta \})$
> $\qquad = \mathrm{Var}_\theta (t(Y_1, ..., Y_n)) + Bias^2((t(.); \theta) + 0$

Cramér-Rao Inequality

$$MSE(\hat{\theta}, \theta) \geq Bias^2(\hat{\theta}, \theta) + \frac{\left(1 + \frac{\partial Bias(\hat{\theta}, \theta)}{\partial \theta} \right)^2}{I(\theta)}$$

> Proof:
> For unbiased estimates: $\theta = \mathrm{E}_\theta(\hat{\theta}) = \int t(y) f(y; \theta) dy$
> $1 = \int t(y) \frac{\partial f(y; \theta)}{\partial \theta} dy$
> $\quad = \int t(y) \frac{\partial \log f(y; \theta)}{\partial \theta} f(y; \theta) dy$
> $\quad = \int t(y) s(y; \theta) f(y; \theta) dy$
> $\quad = \int (t(y) - \theta) (s(\theta; y) - 0) f(y; \theta) dy \qquad$ 1. Bartlett equation $\mathrm{E}_\theta (s(\theta; y)) = 0$
> $\quad = \mathrm{Cov}_\theta (t(Y); s(\theta; Y))$
> $\quad \geq \sqrt{\mathrm{Var}_\theta(t(Y))} \sqrt{\mathrm{Var}_\theta(s(\theta; Y))} \qquad$ Cauchy-Schwarz
> $\quad = \sqrt{MSE(t(Y); \theta)} \sqrt{I(\theta)}$

Kullback-Leibler Divergence   Comparing distributions

$$KL(t, \theta) = \int_{-\infty}^{\infty} \log \frac{f(\tilde{y}; \theta)}{f(\tilde{y}; t)} f(\tilde{y}; \theta) d\tilde{y}$$

The KL divergence is not a distance as it is not symmetric. It is 0 for $t = \theta$ and ≥ 0 otherwise.

> Proof:
> Follows from $\log(x) \leq x - 1 \forall x \geq 0$, with equality for $x = 1$.

$R_{KL}(t(.), \theta)$ is approximated by the MSE.

> Proof:
> $R_{KL}(t(.), \theta) =$
> $= \int_{-\infty}^{\infty} \mathcal{L}_{KL}(t(Y_1, ..., Y_n), \theta) \prod_{i=1}^{n} f(y_i; \theta) dy_i$
> $= \int \int \log \frac{f(\tilde{y}; \theta)}{f(\tilde{y}; t)} f(\tilde{y}; \theta) d\tilde{y} \prod_{i=1}^{n} f(y_i; \theta) dy_i$
> $= \int \int (\log f(\tilde{y}; \theta) - \log f(\tilde{y}; t)) f(\tilde{y}; \theta) d\tilde{y} - \prod_{i=1}^{n} f(y_i; \theta) dy_i$
> $\approx - \int \underbrace{\left(\int \frac{\partial \log f(\tilde{y}; \theta)}{\partial \theta} f(\tilde{y}; \theta) d\tilde{y} \right)}_{0} (t - \theta) \prod_{i=1}^{n} f(y_i; \theta) dy_i$
> $+ \frac{1}{2} \int \underbrace{\left(- \int \frac{\partial^2 \log f(\tilde{y}; \theta)}{\partial \theta^2} f(\tilde{y}; \theta) d\tilde{y} \right)}_{I(\theta)} (t - \theta)^2 \prod_{i=1}^{n} f(y_i; \theta) dy_i$

The last step is approximated by the Taylor Expansion:
$\log f(\tilde{y}, t) \approx \log f(\tilde{y}, \theta) + \frac{\partial \log f(\tilde{y}, \theta)}{\partial \theta} (t - \theta) + \frac{1}{2} \frac{\partial^2 \log f(\tilde{y}, \theta)}{\partial \theta^2} (t - \theta)^2$

3.3   Maximum Likelihood (ML)

Prerequisites

- $Y_i \sim f(y; \theta)$ *i.i.d.*

- $\theta \in \mathbb{R}^p$

- $f(.; \theta)$ Fisher-regular:

  - $\{y : f(y; \theta > 0)\}$ independent of θ

  - Parameter space Θ is open

  - $f(y; \theta)$ twice differentiable

  - $\int \frac{\partial}{\partial \theta} f(y; \theta) dy = \frac{\partial}{\partial \theta} \int f(y; \theta) dy$

Central Functions

- **Likelihood** $L(\theta; y_1, ..., y_n)$: $\prod_{i=1}^n f(y_i; \theta)$

- **log-Likelihood** $l(\theta; y_1, ..y_n)$:
  $\log L(\theta; y_1, ..., y_n) = \sum_{i=1}^n \log f(y_i; \theta)$

- **Score** $s(\theta; y_1, ..., y_n)$: $\frac{\partial l(\theta; y_1, ..y_n)}{\partial \theta}$

- **Fisher-Information** $I(\theta)$: $-E_\theta \left(\frac{\partial s(\theta; Y)}{\partial \theta} \right)$

- **observed Fisher-Information** $I_{obs}(\theta)$: $-E_\theta \left(\frac{\partial s(\theta; y)}{\partial \theta} \right)$

Attributes of the Score-Function

first Bartlett-Equation:

$$E\left(s(\theta; Y) \right) = 0$$

> Proof:
> $$1 = \int f(y; \theta) dy$$
> $$0 = \frac{\partial 1}{\partial \theta} = \int \frac{\partial f(y; \theta)}{\partial \theta} dy = \int \frac{\partial f(y; \theta)/\partial \theta}{f(y; \theta)} f(y; \theta) dy$$
> $$= \int \frac{\partial}{\partial \theta} \log f(y; \theta) f(y; \theta) dy = \int s(\theta; y) f(y; \theta) dy$$

second Bartlett-Equation:

$$\text{Var}_\theta \left(s(Y; \theta) \right) = E_\theta \left(-\frac{\partial^2 \log f(Y; \theta)}{\partial \theta^2} \right) = I(\theta)$$

> Proof:
> $$0 = \frac{\partial 0}{\partial \theta} = \frac{\partial}{\partial \theta} \int \frac{\partial}{\partial \theta} \log f(y; \theta) f(y; \theta) dy \qquad \text{see above}$$
> $$= \int \left(\frac{\partial^2}{\partial \theta^2} \log f(y; \theta) \right) f(y; \theta) dy$$
> $$+ \int \frac{\partial \log f(y; \theta)}{\partial \theta} \frac{\partial f(y; \theta)}{\partial \theta} dy$$
> $$= E_\theta \left(\frac{\partial^2}{\partial \theta^2} \log f(Y; \theta) \right)$$
> $$+ \int \frac{\partial \log f(y; \theta)}{\partial \theta} \frac{\partial \log f(y; \theta)}{\partial \theta} f(y; \theta) dy$$
> $$\Leftrightarrow E_\theta \left(s(\theta; Y) s(\theta; Y) \right) = E_\theta \left(-\frac{\partial^2}{\partial \theta^2} \log f(Y; \theta) \right)$$
> Bartlett's second equation holds then as $E\left(s(\theta; Y) \right) = 0$

ML-Estimate

$$\hat{\theta}_{ML} = \arg \max l(\theta; y_1, ...y_n)$$

for Fisher-regular distributions: $\hat{\theta}_{ML}$ has aysmptotically the smallest variance, given by the Cramér-Rao inequality,
$s\left(\hat{\theta}_{ML}; y_1, ..., y_n \right) = 0$
$\hat{\theta} \overset{a}{\sim} N\left(\theta, I^{-1}(\theta) \right)$
If the true model is unknown, the distribution is
$\hat{\theta} \overset{a}{\sim} N\left(\theta, I^{-1}(\theta) V(\theta) I^{-1}(\theta) \right)$ with $V(\theta)$ variance of the score function.
The ML-estimate is invariant: $\hat{\gamma} = g(\hat{\theta})$ if $\gamma = g(\theta)$.

> Proof:
> $\gamma = g(\theta) \Leftrightarrow \theta = g^{-1}(\gamma)$
> For the log-likelihood of γ at the location $\hat{\theta}$ holds:
> $$\frac{\partial l(g^{-1}(\hat{\gamma}))}{\partial \gamma} = \frac{\partial g^{-1}(\gamma)}{\partial \gamma} \underbrace{\frac{\partial l(\hat{\theta})}{\partial \theta}}_{=0} = 0$$

Then, the Fisher information is $\frac{\partial \theta}{\partial \gamma} I(\theta) \frac{\partial \theta}{\partial \gamma}$

> Proof:
> $$I_\gamma(\gamma) = -E\left(\frac{\partial^2 l(g^{-1}(\hat{\gamma}))}{\partial \gamma^2} \right) = -E\left(\frac{\partial}{\partial \gamma} \left(\frac{\partial g^{-1}(\gamma)}{\partial \gamma} \frac{\partial l(\theta)}{\partial \theta} \right) \right)$$
> $$= -E\left(\underbrace{\frac{\partial^2 g^{-1}(\gamma)}{\partial \gamma} \frac{\partial l(\theta)}{\partial \theta}}_{\text{Erwartungswert } 0} + \frac{\partial g^{-1}(\gamma)}{\partial \gamma} \frac{\partial^2 l(\theta)}{\partial \theta^2} \frac{\partial g^{-1}(\gamma)}{\partial \gamma} \right)$$
> $$= \frac{\partial g^{-1}(\gamma)}{\partial \gamma} I(\theta) \frac{\partial g^{-1}(\gamma)}{\partial \gamma} = \frac{\partial \theta}{\partial \gamma} I(\theta) \frac{\partial \theta}{\partial \gamma}$$

Delta rule: $\gamma \overset{a}{\sim} N(\hat{\gamma}, \frac{\partial \theta}{\partial \gamma} I^{-1}(\theta) \frac{\partial \theta}{\partial \gamma})$

Numerical computation of the ML estimate   Fisher-Scoring as statistical version of the Newton-Raphson procedure

1. Initialize $\theta_{(0)}$

2. Repeat: $\theta_{(t+1)} := \theta_{(t)} + I^{-1}(\theta_{(t)}) s(\theta_{(t)}; y)$

3. Stop if $\|\theta_{(t+1)} - \theta_{(t)}\| < \tau$; return $\hat{\theta}_{ML} = \theta_{(t+1)}$

> Proof:
> $$0 = s(\hat{\theta}_{ML}; y) \overset{Taylor}{\underset{Series}{\approx}} s(\theta; y) + \frac{\partial s(\theta; y)}{\partial \theta} (\hat{\theta}_{ML} - \theta) \Leftrightarrow$$
> $$\hat{\theta}_{ML} \approx \theta - \left(\frac{\partial s(\theta; y)}{\partial \theta} \right)^{-1} s(\theta; y) \approx \theta - I^{-1}(\theta) s(\theta; y)$$
> As $\frac{\partial s(\theta; y)}{\partial \theta}$ is often complicated, its expectation $I(\theta)$ is used.

The second part in 2 can be weighted with a step size δ or $\delta(t)$ $\in (0, 1)$, e.g. to ensure convergence.
If $I(\theta)$ can't be analytically derived, simulation from $f(y; \theta_{(t)})$ can be used. For the exponential family, step 2 then changes to $\theta_{(t+1)} := \theta_{(t)} + \hat{\text{Var}}_{\theta_{(t)}}(t(Y))^{-1} E_{\theta_{(t)}}(t(Y))$ as the ML estimate is the expectation.

Log Likelihood Ratio

$$lr(\theta, \hat{\theta}) := l(\hat{\theta}) - l(\theta) = \log \frac{L(\hat{\theta})}{L(\theta)}$$

with $2 \cdot lr(\theta, \hat{\theta}) \overset{a}{\sim} \chi_1^2$

Proof:

$$l(\theta) \overset{Taylor}{\underset{Series}{\approx}} l(\hat{\theta}) + \frac{\partial l(\hat{\theta})}{\partial \theta}(\theta - \hat{\theta}) + \frac{1}{2} \underbrace{\frac{\partial^2 l(\hat{\theta})}{\partial \theta^2}}_{\approx I^{-1}(\theta) s(\theta; Y)} (\underbrace{\theta - \hat{\theta}}_{\approx -I(\theta)})^2$$

with $\frac{\partial l(\hat{\theta})}{\partial \theta} = 0$

$$\approx l(\hat{\theta}) - \frac{1}{2} \frac{s^2(\theta, Y)}{I(\theta)}$$

$s(\theta, Y)$ is asymptotically normal.

If $\theta \in \mathbb{R}^p$ the corresponding distribution is χ_p^2.

Relation to Kullback-Leibler divergence

$$\hat{\theta}_{ML} = \arg \min KL(g, f)$$

with f distributional model used and g true model

Proof:

$$KL(g, f) = \int \log \frac{g(y)}{f(y)} g(y) dy$$

$$= \int \log(g(y)) g(y) dy - \int \log(f(y)) g(y) dy$$

To minimize that, the second component needs to be maximized. Its derivative is $\int s(\theta; y) g(y) dy = E_g(s(\theta; Y)) = 0$

3.4 Consistency and Sufficiency

(Weak) Consistency

$$MSE(\hat{\theta}, \theta) \overset{n \to \infty}{\longrightarrow} 0 \Rightarrow \hat{\theta} \text{ consistent}$$

Proof:
$P(|\hat{\theta} - E_{\hat{\theta}}| \geq \delta) \leq \frac{Var_\theta(\hat{\theta})}{\delta^2}$ using the inequality of Chebyshev
and $MSE(t(.), \theta) = Var_\theta(t(Y_1, ..., Y_n)) + Bias^2((t(.); \theta))$

Statistic

$$t : \mathbb{R}^n \to \mathbb{R}$$

$t(Y_1, ..., Y_n)$ depends on sample size n and is a random variable

Suffizienz

A statistic $t(y_1, ..., y_n)$ is sufficient for θ, if the conditional distribution $f(y_1, ..., y_n | t_0 = t(y_1, ..., y_n); \theta)$ is independent of θ.

Neyman criterion:

$t(Y_1, ..., Y_n)$ suffizient $\Leftrightarrow f(y; \theta) = h(y) g(t(y); \theta)$

Proof:
"⇒":

$$f(y; \theta) = \underbrace{f(y | t = t(y); \theta)}_{h(y)} \underbrace{f_t(t | y; \theta)}_{g(t(y); \theta)}$$

"⇐":

$$f_t(t; \theta) = \int_{t=t(y)} f(y; \theta) dy = \int_{t=t(y)} h(y) g(t; \theta) dy$$

Therefore:

$$f(y | t = t(y); \theta) = \frac{f(y, t = t(y); \theta)}{f_t(t, \theta)} = \begin{cases} \frac{h(y) g(t; \theta)}{g(t; \theta)} & t = t(y) \\ 0 & \text{otherwise} \end{cases}$$

Minimal Sufficiency:

$t(.)$ is sufficient and $\forall \tilde{t}(.) \exists h(.)$ s.t. $t(y) = h(\tilde{t}(y))$

4 Statistical Hypothesis Testing

4.1 Significance and Confidence Intervals

Significance Test

Assuming two states H_0 and H_1 and two corresponding decisions "H_0" and "H_1", a decision rule (a threshold $c \in \mathbb{R}$ for the test statistic $T(X)$) is constructed s.t.:

$$P(\text{``}H_1\text{''}|H_0) \leq \alpha$$

	"H_0"	"H_1"
H_0	$1-\alpha$ (correct)	α (type I error)
H_1	β (type II error)	$1-\beta$ (correct)

Power   concerns the type II error

$$power = P(\text{``}H_1\text{''}|H_1) = 1 - \beta$$

p-Value   measures the amount of evidence against H_0

$$p - value \leq \alpha \Leftrightarrow \text{``}H_0\text{''}$$

Confidence Interval

$$[t_l(Y), t_r(Y)] \text{ Confidence Interval}$$
$$\Leftrightarrow$$
$$P_\theta\left((t_l(Y) \leq \theta \leq t_r(Y)\right) \geq 1-\alpha$$

with $1 - \alpha$ confidence level und α significance level

Corresponding Test

$$\theta \notin [t_l(y), t_r(y)] \Leftrightarrow \text{``}H_1\text{''}$$

Specificity   or True Negative Rate ($1-$empirical type I error)

$$TNR = \frac{\#TN}{\#N} = \frac{\#TN}{\#TN + \#FP}$$

Sensitivity   or True Positive Rate, Recall (empirical power)

$$TPR = \frac{\#TP}{\#P} = \frac{\#TP}{\#TP + \#FN}$$

4.2 Tests for One Sample

Normal Distribution $X_i \overset{iid}{\sim} N(\mu, \sigma^2)$

Test for μ, known σ^2 (Simple Gauss-Test)

$H_0\colon \mu = \mu_0 \quad vs. \quad H_1\colon \mu \neq \mu_0$

$$T(X) = \frac{\bar{X} - \mu_0}{\sigma} \overset{H_0}{\sim} N(0,1)$$

Test for μ, unknown σ^2 (Simple t-Test)

$H_0\colon \mu = \mu_0 \quad vs. \quad H_1\colon \mu \neq \mu_0$

$$T(X) = \frac{\bar{X} - \mu_0}{\hat{\sigma}/\sqrt{n}} \overset{H_0}{\sim} t_{n-1}$$

with $\hat{\sigma} = \sqrt{\frac{1}{n-1} \sum_{i=1}^n (Y_i - \bar{Y})^2}$

ML Estimate $\hat{\theta} \overset{a}{\sim} N(\theta, I^{-1}(\theta))$

Wald Test

$H_0\colon \theta = \theta_0 \quad vs. \quad H_1\colon \theta \neq \theta_0$

$$T(X) = |\hat{\theta} - \theta_0| \overset{H_0}{\sim} N(0, I^{-1}(\theta_0))$$

As $\hat{\theta}$ converges to θ_0 under H_0, it can also be used to calculate the variance: $I^{-1}(\hat{\theta})$.

Score Test

$H_0\colon \theta = \theta_0 \quad vs. \quad H_1\colon \theta \neq \theta_0$

$$T(X) = |s(\theta_0; y)| \overset{H_0}{\sim} N(0, I(\theta_0))$$

Advantage compared to the Wald Test: $\hat{\theta}$ does not have to be calculated.

Likelihood Ratio Test

$H_0\colon \theta = \theta_0 \quad vs. \quad H_1\colon \theta \neq \theta_0$

$$T(X) = 2(l(\hat{\theta}) - l(\theta)) \overset{H_0}{\sim} \chi_1^2$$

Neyman-Pearson Test

$H_0\colon \theta = \theta_0 \quad vs. \quad H_1\colon \theta = \theta_1$

$$T(X) = l(\theta_0) - l(\theta_1)$$

For a given significance level α, the Neyman Pearson Test is the most powerful test for comparing two estimates for θ.

Proof:

Decision rule of the NP-Test: $\varphi^* = \begin{cases} 1 & if \frac{f(y;\theta_0)}{f(y;\theta_1)} \leq e^c \\ 0 & \text{otherwise} \end{cases}$

Need to show: $P(\varphi(Y)=1|\theta_1) \leq P(\varphi^*(Y)=1|\theta_1) \; \forall \varphi$

$P(\varphi^*=1|\theta_1) - P(\varphi=1|\theta_1) =$

$$= \int \{\varphi^*(y) - \varphi(y)\} f(y;\theta_1) dy$$

$$\geq \frac{1}{e^c} \int_{\varphi^*=1} \{\varphi^*(y) - \varphi(y)\} f(y;\theta_0) dy \quad f(y;\theta_1) \geq \frac{f(y;\theta_0)}{e^c}$$

$$+ \frac{1}{e^c} \int_{\varphi^*=0} \{\varphi^*(y) - \varphi(y)\} f(y;\theta_0) dy \quad f(y;\theta_1) \leq \frac{f(y;\theta_0)}{e^c}$$

$$= \frac{1}{e^c} \int \{\varphi^*(y) - \varphi(y)\} f(y;\theta_0) dy = 0$$

As $\alpha = \int \varphi^*(y) f(y;\theta_0) dy = \int \varphi(y) f(y;\theta_0) dy$

4.3 Tests for Goodness of Fit

Discrete (Chi-Squared)
$H_0 \colon X_i \sim F_0$ $\;vs.\;$ $H_1 \colon X_i \sim F \neq F_0$

$$T(X) = \sum_{k=1}^{K} \frac{(n_k - l_k)^2}{l_k} \overset{H_0}{\sim} \chi^2_{K-1-p}$$

with the following contingency table:

	1	2		K
observed	n_1	n_2	...	n_K
expected under H_0	l_1	l_2	...	l_K

$l_k > 5$ and $l_k > n - 5$ for the χ^2_{K-1-p}-distribution to hold, F_0 needs to be known, but its p parameters can be estimated. The test can be applied to discretized continuous variables.

Continuous (Kolmogorov-Smirnov Test)
$H_0 \colon X_i \sim F_0$ $\;vs.\;$ $H_1 \colon X_i \sim F \neq F_0$

$$T(X) = \sup_x |F_n(x) - F(x;\theta)| \overset{H_0}{\sim} KS$$

with the distribution function $F(x;\theta)$ and the empirical counterpart $F_n(x) = \frac{1}{n} \sum_{i=1}^{n} \mathbb{1}_{\{X_i \leq x\}}$

> Proof:
>
> $P(\sup_x |F_n(x) - F(x;\theta)| \leq t) =$
>
> $= P(\sup_y |F^{-1}(y;\theta) - x| \leq t)$ $\qquad \substack{x \in [0,1],\, x = F^{-1}(y;\theta) \\ F(F^{-1}(y;\theta);\theta) = y}$
>
> $\overset{*}{=} P(\sup_y |\frac{1}{n} \sum_{i=1} \mathbb{1}_{\{U_i \leq y\}} - y| \leq t)$ $\quad$ with $U_i \sim U(0,1)$
>
> $^* F_n(F^{-1}(y;\theta)) = \frac{1}{n} \sum_{i=1}^{n} \mathbb{1}_{\{X_i \leq F^{-1}(y;\theta)\}} = \frac{1}{n} \sum_{i=1}^{n} \mathbb{1}_{\{F(y;\theta) \leq y\}}$

For an estimated parameter the distribution of $T(X)$ is not independent of F_0: $T(X) \overset{H_0}{\sim} KS$ only holds asymptotically.

Pivotal Statistic

$$g(Y;\theta) \text{ pivotal}$$
$$\Leftrightarrow$$
$$\text{Distribution of } g(Y;\theta) \text{ independent of } \theta$$

Approximative Pivotal Statistic
$H_0 \colon X_i \sim F$ pivotal $\;vs.\;$ $H_1 \colon X_i \sim F$ not pivotal

$$g(\hat{\theta};\theta) = \frac{\hat{\theta} - \theta}{\sqrt{\mathrm{Var}(\hat{\theta})}} \overset{\alpha}{\sim} \mathrm{N}(0,1)$$

with $\hat{\theta} = t(Y) \overset{\alpha}{\sim} \mathrm{N}(\theta, \mathrm{Var}(\hat{\theta}))$

$$KI = \left[\hat{\theta} - z_{1-\frac{\alpha}{2}} \sqrt{\mathrm{Var}(\hat{\theta})}, \hat{\theta} + z_{1-\frac{\alpha}{2}} \sqrt{\mathrm{Var}(\hat{\theta})} \right]$$

> Proof:
> $1 - \alpha \approx P\left(z_{\frac{\alpha}{2}} \leq \frac{\hat{\theta} - \theta}{\sqrt{\mathrm{Var}(\hat{\theta})}} \leq z_{1-\frac{\alpha}{2}} \right)$

4.4 Multiple Tests

Family-Wise Error Rate (FWER) $\quad$ as p-$value \sim U(0,1)$

For m tests:

$$\alpha \leq P\left(\cup_{k=1}^{m} (p_k \leq \alpha) | H_{0k}, k = 1, ..., m \right) \leq m\alpha$$

$$FWER := P(\exists k : \text{``}H_1 k\text{''} | \forall k : H_0 k)$$

Bonferoni Adjustment

$$\alpha_B = \frac{\alpha}{m}$$

Šidák Adjustment $\quad$ only for independent tests

$$\alpha_S = 1 - (1 - \alpha)^{1/m}$$

> Proof:
> $\alpha \overset{!}{=} P\left(\cup_{k=1}^{m} (p_k \leq \alpha) | H_{0k}, k = 1, ..., m \right)$
> $= 1 - (1 - \alpha)^{1/m}$

Holm's Procedure $\quad$ also takes power into account
Order the p-values: $p_{(1)} \leq ... \leq p_{(m)}$
Step $x \in \mathbb{N}^+$: if $p(x) > \frac{\alpha}{m+1-x}$ reject H_{01} to H_{0x} and stop, else move on to step $x + 1$.

False Discovery Rate (FDR) $\quad$ balances type I and II errors, especially for $n << m$ problems

$$FDR = \mathrm{E}\left(\frac{\#\text{``}H1\text{''} | H_0}{\#\text{``}H1\text{''}} \right)$$

Order the p-values: $p_{(1)} \leq ... \leq p_{(m)}$, choose $\alpha \in (0,1)$
j is largest index s.t. $p(j) \leq \alpha j/m$, reject all $H_0 i$ for $i \leq j$

It can be shown that $FDR \leq m_0 \alpha / m$, with $m_0 = \#H_0$

5 Regression

5.1 Models

5.1.1 Simple Linear Model

Theoretical Model

$$y_i = \beta_0 + \beta_1 x_i + u_i$$

Empirical Model

$$y_i = \hat{\beta}_0 + \hat{\beta}_1 x_i + e_i$$

$$\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i$$

Assumptions

- **Independent Observations** $y_1, ... y_n$ are independent

- **Linearity of the Mean** $E(Y|x) = \beta_0 + \beta_1 x$ or $\mathrm{E}(e|x) = 0$

- **Constant Variation** $Var(Y|x) = \sigma^2$

For the normal linear model:

- **Normality** $e|x \sim \mathrm{N}(0, \sigma^2)$; $Y|x \sim \mathrm{N}(\hat{y}, \sigma^2)$

Attributes of the Regression Line

$$\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i = \bar{y} + \hat{\beta}_1(x_i - \bar{x})$$
$$\hat{e}_i = y_i - \hat{y}_i = y_i - (\hat{\beta}_0 + \hat{\beta}_1 x_i)$$
$$= y_i - (\bar{y} + \hat{\beta}_1(x_i - \bar{x}))$$
$$\sum_{i=1}^{n} \hat{e}_i = \sum_{i=1}^{n} y_i - \sum_{i=1}^{n} \bar{y} - \hat{\beta}_1 \sum_{i=1}^{n}(x_i - \bar{x})$$
$$= n\bar{y} - n\bar{y} - \hat{\beta}_1(n\bar{x} - n\bar{x}) = 0$$
$$\bar{\hat{y}} = \frac{1}{n}\sum_{i=1}^{n} \hat{y}_i = \frac{1}{n}(n\bar{y} + \hat{\beta}_1(n\bar{x} - n\bar{x})) = \bar{y}$$

Estimates (OLS)

$$\hat{\beta}_1 = \frac{Cov(x,y)}{Var(x)} = \frac{S_{xy}}{S_{xx}} = \frac{S_{xy}}{\sqrt{S_{xx}S_{yy}}} \cdot \sqrt{\frac{S_{yy}}{S_{xx}}} = r\sqrt{\frac{S_{yy}}{S_{xx}}}$$

> Proof:
> $$Cov(x,y) = Cov(x, \hat{\beta}_0 = \bar{y} - \hat{\beta}_1\bar{x}) = \hat{\beta}_1 Var(x)$$
> $$\iff \hat{\beta}_1 = \frac{Cov(x,y)}{Var(x)}$$

$$\hat{\beta}_0 = \bar{y} - \hat{\beta}_1\bar{x}$$

> Proof:
> $$E[y] = E\left[\hat{\beta}_0 + \hat{\beta}_1 x + \hat{e}\right] \iff \hat{\beta}_0 = E[y] - \hat{\beta}_1 E[x]$$

The estimates are the same as for the ML procedure.

Estimates (ML) $Y|x \sim \mathrm{N}(\beta_0 + \beta_1 x, \sigma^2)$

$$\hat{\beta}_0 = \frac{1}{n}\sum_{i=1}^{n} y_i - \frac{1}{n}\sum_{i=1}^{n} x_i \hat{\beta}_1$$
$$\hat{\beta}_1 = \sum_{i=1}^{n} x_i(y_i - \hat{\beta}_0 / \sum_{i=1}^{n} x_i^2$$
$$\hat{\sigma}^2 = \frac{1}{n}\sum_{i=1}^{n}(y_i - \hat{\beta}_0 - x_i\hat{\beta}_1)^2$$

The β-estimates are the same as for the OLS procedure.

> Proof:
> $$l(\beta_0, \beta_1, \sigma^2) = \sum_{i=1}^{n}\left\{-\frac{1}{2}\sigma^2 - \frac{1}{2}\frac{(y_i - \beta_0 - \beta_1 x_i)^2}{\sigma^2}\right\}$$

5.1.2 Multivariate Linear Model

Theoretical Model

$$Y = X\beta + u$$

Empirical Model

$$Y = X\hat{\beta} + e$$
$$\hat{Y} = X\hat{\beta}$$

$$y = (y_1, ..., y_n)^T, e = (e_1, ..., e_n)^T, X = \begin{pmatrix} 1 & x_{11} & ... & x_{1p} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & x_{n1} & ... & x_{np} \end{pmatrix}$$

Assumptions

- **Independent Observations** $y_1, ... y_n$ are independent

- **Linearity of the Mean** $\mathrm{E}(Y|x_{1:p}) = X\beta$ or $\mathrm{E}(e|x_{1:p}) = 0$

- **Constant Variation** $Var(Y|x) = \sigma^2$

For the normal linear model:

- **Normality** $e_i|x_{1:p} \sim \mathrm{N}(0, \sigma^2)$; $Y|x \sim \mathrm{N}(\hat{y}, \sigma^2)$

Estimates (ML) $Y|x_{1:p} \sim \mathrm{N}(X\beta, \sigma^2)$

$$\hat{\beta} = \left(X^T X)^{-1}\right) X^T y$$
$$Var(\hat{\beta}) = \sigma^2(X^T X)^{-1} = I^{-1}(\beta)$$

> Proof:
> $$l(\beta, \sigma^2) = -\frac{n}{2}\log\sigma^2 - \frac{1}{2\sigma^2}(y - X\beta)^T(y - X\beta)$$

The estimates are the same as for the OLS procedure.

β is the **B**est **L**inear **U**nbiased **E**stimator

Proof:
Unbiased because of the Gauß-Markov Theorem: $\mathrm{E}(\hat{\beta}) = (X^TX)^{-1}X^T\mathrm{E}(Y|X) = (X^TX)^{-1}X^TX\beta = \beta$

Proof:
$H := X(X^TX)^{-1}X^T$ hat matrix; $HH = H = H^T$ (idempotent)
$$\mathrm{E}((Y-X\hat{\beta})^T(Y-X\hat{\beta})) = \mathrm{E}((Y^T(I_n-H)^T((I_n-H)Y))$$
$$= \mathrm{E}(tr(Y^T(I_n-H)Y))$$
$$= \mathrm{E}(tr((I_n-H)YY^T))$$
$$= tr((I_n-H)\mathrm{E}(YY^T))$$
$$= tr((I_n-H)\mathrm{E}(X\beta\beta^TX^T + \sigma I_n))$$
$$= \sigma^2 tr((I_n-H))$$
$$= \sigma^2(n-p)$$

$$\hat{\sigma}^2 = \frac{1}{n}(y-X\hat{\beta})^T(y-X\hat{\beta}); \quad \hat{\beta} \sim \mathrm{N}(\beta, \sigma^2(X^TX)^{-1})$$

The ML-estimate for σ^2 is biased.

$$s^2 = \frac{1}{n-p}(y-X\hat{\beta})^T(y-X\hat{\beta}); \quad \hat{\beta} \sim t_{n-p}(\beta, s^2(X^TX)^{-1})$$
with s an unbiased estimator

5.1.3 Bayesian Linear Model

Prior flat prior

$$f_{\beta,\sigma^2}(\beta, \sigma^2) = \frac{1}{\sigma^2}$$

Posterior
Resulting posterior:
$$f_{post}(\beta, \sigma^2|y) \propto (\sigma^2)^{-\frac{n}{2}+1}e^{-\frac{1}{2\sigma^2}(y-X\beta)^T(y-X\beta)}$$
Note: $f_{post}(\beta, \sigma^2|y) = f(\beta|\sigma^2, y)f(\sigma^2|y)$

$$\beta|\sigma^2, y \sim \mathrm{N}\left(\hat{\beta}, \sigma^2(X^TX)^{-1}\right)$$
$$\sigma^2|y \sim \mathrm{IG}\left(\frac{n-p}{2}, \frac{s^2(n-p)}{2}\right)$$
$$\beta|y \sim t_{n-p}\left(\hat{\beta}, s^2(X^TX)^{-1}\right)$$

The two distributions for β mirror the results for $\hat{\beta}$ in the linear model.

5.1.4 Quantile Regression

Prediction Interval range of $1 - \alpha$ fraction of the data

$$Var(\hat{Y}|x_{1:p}) = Var(X\hat{\beta}) + \sigma^2$$

Determined by estimation variance (usually captured by confidence intervals) plus residual variance.

Quantile
$$Q(\tau) = \inf\{y : F(y) \geq \tau\}$$
If F is invertable: $Q(\tau) = F^{-1}(\tau)$, $\tau \in (0,1)$

Model
$$Q(\tau|x_{1:p}) = X\beta$$
For median regression: $\hat{\beta} = \arg\ \min \sum_{i=1}^{n} |y_i - x_i^T\beta|$
In general:
$$\hat{Q}(\tau) = \arg\min_{\beta}\left(\sum_{i=1}^{n}\delta_\tau(y_i - x_i^T\beta)\right)$$

with check function $\delta_\tau(y) = y\left(\tau - \mathbb{1}_{\{y<0\}}\right)$

Proof:
$$Q(\tau) = \arg\min_{q}\mathrm{E}(\delta_\tau(Y-q))$$
$$= \arg\min_{q}\left\{(\tau-1)\int_{-\infty}^{q}(y-q)f(y)dy + \tau\int_{q}^{\infty}(y-q)f(y)dy\right\}$$
Differentiating w.r.t. q gives $(\tau-1)\int_{-\infty}^{q}f(y)dy - \tau\int_{q}^{\infty}f(y)dy$
$= (1-\tau)F(q) - \tau(1-F(q)) = F(q) - \tau$

Estimates
The estimates for β can be computed with linear programming and are normally distributed with mean β.

5.1.5 Flexible Regression

Assumptions

- **Independent Observations** $y_1, ... y_n$ are independent

- **Constant Variation** $Var(Y|x) = \sigma^2$

- **Normality** $e_i|x_{1:p} \sim \mathrm{N}(0, \sigma^2)$; $Y|x \sim \mathrm{N}(\hat{y}, \sigma^2)$

Knot Placement

- equidistant

- based on quantiles (more structure where data is dense)

- all data points plus penalization

Penalized Regression Splines

$$||y - X\beta||^2 + \lambda \int_{x_1}^{x_n} \left[f''(x) \right]^2 dx = ||y - X\beta||^2 \lambda \beta^T D\beta$$

$$l_p(\beta, \sigma^2, \lambda) = l(\beta, \sigma^2) - \frac{\lambda}{2\sigma_\epsilon^2} \beta^T D\beta$$

$$\hat{\beta} = (X^T X + \lambda D)^{-1} X^T y$$

Difference Penalty

- first order: $\beta^T D\beta = \sum_{j=1}^{p} (\beta_{j+1} - \beta_j)^2$

- second order: $\beta^T D\beta = \sum_{j=1}^{p} (\beta_{j+1} - 2\beta_j + \beta_{j-1})^2$

Choosing λ   Model complexity

$$\dim(\lambda) = tr\left\{ (X^T X + \lambda D)^{-1}(X^T X) \right\}$$

$$AIC(\lambda) = fit(\lambda) + 2\dim(\lambda)$$

Numerically complex. Alternative: **Bayes**

$$\beta \sim N(0, \sigma_\beta^2 D^-) \text{ with } (D^-)^- = D \text{ (generalized inverse)}$$

$$\log f(\beta, \sigma^2; \sigma_\beta^2 | y) \propto l(\beta, \sigma^2) - \frac{rk(D^-)}{2} \log(\sigma_\beta^2) - \frac{1}{2\sigma_\beta^2} \beta^T D^- \beta$$

As $\lambda = \frac{1}{\sigma_\beta^2}$, marginal posterior for σ_β^2 can be derived. E. g. set λ to the posterior mode estimate.

5.1.6   Generalized Regression

Assumptions

- **Independent Observations** $y_1, ... y_n$ are independent

- **Linearity of the Mean** $E(Y | x_{1:p}) = X\beta$ or $E(e | x_{1:p}) = 0$

- **Exponential Family** $Y | x \sim \exp\left\{ t(y)\theta(x) - \kappa(\theta(x)) \right\} h(y)$

Link Function
Linear predictor $\eta = X\beta$; $\mu = \frac{\partial \kappa(\theta)}{\partial \theta} = E(t(Y); \theta)$

$$\mu = g^{-1}(\eta)$$

If $\lambda = 0$, *canonical link*:

$$g(\theta) = \eta$$

- score function: $s(\beta) = X^T \left(t(y) - E(t(Y); \eta) \right)$

- estimate $\hat{\beta} = X^T E(t(Y); \hat{\eta}) = X^T t(y)$

- Fisher matrix $I(\beta) = X^T W X$
  with W diagonal and $W_{ii} = \frac{\partial^2 \kappa(\eta_i)}{\partial \eta^2} = Var(t(Y_i), \eta_i)$

Examples:

- **Logistic**: $\text{logit} P(Y_i = 1 | x_i) = \log \frac{P(Y_i = 1 | x_i)}{1 - P(Y_i = 1 | x_i)} = \eta$
  $Var(Y_i | x_i) = P(Y_i = 1 | x_i) \cdot (1 - P(Y_i = 1 | x_i))$

- **Poisson**: $log E(Y_i | x_i) = \eta$
  $Var(Y_i | x_i) = E(Y_i | x_i) = e^\eta$

5.1.7   Weighted Regression

Different Precision   variance heterogeneity: $e_i \sim N(0, \sigma_i^2)$

$$l(\beta, \sigma^2) = -\frac{n}{2} \log \sigma^2 - \frac{1}{2\sigma^2 (y - X\beta)^T (y - X\beta)}$$

with $W = diag(\frac{1}{a_1}, ..., \frac{1}{a_n})$ and $a_i = \frac{\sigma_i^2}{\sigma^2}$

$$\hat{\beta}_{ML} = (X^T W X)^{-1}(X^T W y)$$

$$Var(\hat{\beta}_{ML}) = \sigma^2 (X^T W X)^{-1}$$

Different Group Representation

$$Y_i | x_{i,1:p}, z_i \sim N(x_{i,1:p} \beta_{z_i}, \sigma^2)$$

with z_i indicating group affiliation

5.2   Analysis of Variances (ANOVA)

$$SS_{Total} = SS_{Explained} + SS_{Residual}$$

with

$$SS_{Total} = \sum_{i=1}^{n} (y_i - \bar{y})^2$$

$$SS_{Explained} = \sum_{i=1}^{n} (\hat{y}_i - \bar{y})^2$$

$$SS_{Residual} = \sum_{i=1}^{n} (y_i - \hat{y}_i)^2 = \sum_{i=1}^{n} e_i^2 = S_{yy} - \hat{\beta}^2 S_{xx}$$

5.3   Goodness of Fit

5.3.1 Coefficient of Determination

$$R^2 = \frac{SS_{Explained}}{SS_{Total}} = 1 - \frac{SS_{Residual}}{SS_{Total}} = r^2 \qquad \qquad \text{Range: } 0 \leq R^2 \leq 1$$

6 Bayesian Statistics

6.1 Basics

Bayes Theorem
$$P(A|B) = \frac{P(B|A)P(A)}{P(B)} \qquad \text{für } P(A), P(B) > 0$$

or more general:
$$f_{post}(\theta|X) = \frac{f(X|\theta) \cdot f_\theta(\theta)}{\int f(X|\tilde{\theta})f_\theta(\tilde{\theta})d\tilde{\theta}}$$
$$= C \cdot f(X|\theta) \cdot f_\theta(\theta) \quad \text{choose C so that } \int f(\theta|X) = 1$$
$$\propto f(X|\theta) \cdot f_\theta(\theta)$$

Point Estimates
$$\hat{\theta}_{postmean} = E_0(\vartheta|x) = \int_{\vartheta \in \Theta} \vartheta f_\theta(\vartheta|x)d\vartheta$$
$$\hat{\theta}_{postmode} = \arg\max_\vartheta f_\theta(\vartheta, x)$$
$$\hat{\theta}_{Bayesrisk} = \arg\min_{t(.)} R_{Bayes}(t(.))$$

with Bayes risk: $R_{Bayes}(t(.)) = \int_\Theta R(t(.), \vartheta) f_\theta(\vartheta)d\vartheta$
$$\hat{\theta}_{postBayesrisk} = \arg\min_{t(.)} R_{postBayes}(t(.)|y)$$

with posterior Bayes risk:
$$R_{postBayes}(t(.)|y) = \int L(t(y), \vartheta) f_\theta(\vartheta|y) = E_{\theta|y}(L(t(y), \theta)|y$$

Credibility Interval
$$P_\theta(\theta \in [t_l(y), t_r(y)]\,|y) = \int_{t_l(y)}^{t_r(y)} f_\theta(\vartheta|y)d\vartheta = 1 - \alpha$$

- symmetric: $\int_{-\infty}^{t_l(y)} f_\theta(\vartheta|y)d\vartheta = \int_{t_r(y)}^{\infty} f_\theta(\vartheta|y)d\vartheta = \frac{\alpha}{2}$

- highest density: $HDI = \theta|f_\theta(\theta|y) \geq c$, choose c s.t.
$\int_{\vartheta \in HDI(y)} f_\theta(\vartheta|y)d\vartheta = 1 - \alpha$

Bayes Factor evidence contained in data for M_1 vs. M_2
$$\frac{P(M_1|y)}{P(M_0|y)} = \underbrace{\frac{f(y|M_1)}{f(y|M_0)}}_{\text{Bayes Factor}} \frac{P(M_1)}{P(M_0)}$$

with marginal likelihood $f(y|M_i) = \int f(y|\vartheta)f_\theta(\vartheta|M_i)d\vartheta$

Priors

Flat (uninformative) Prior

$f_\theta(\theta) = const.$ for $\theta > 0$, therefore: $f(\theta|X) = C \cdot f(X|\theta)$
As $\int f_\theta(\theta) = 1$ not possible like this, this is not a real density.
Changes for transformations of the parameter.

> Proof: For $\gamma = g(\theta)$: $f_\gamma(\gamma) = f_\theta(g^{-1}(\gamma)) \left| \frac{\partial g^{-1}(\gamma)}{\partial \gamma} \right|$

No prior is truly uninformative.

Jeffrey's Prior
Remains unchanged for transformations of the parameter.
For Fisher-regular distributions: $f(\theta) \propto \sqrt{I_\theta(\theta)}$

> Proof:
> For $\gamma = g(\theta)$ and $f_\theta(\theta) = \sqrt{I_\theta(\theta)}$:
> $$f_\gamma(\gamma) \propto f_\theta(g^{-1}(\gamma)) \left| \frac{\partial g^{-1}(\gamma)}{\partial \gamma} \right| \propto \sqrt{\frac{\partial g^{-1}(\gamma)}{\partial \gamma} I_\theta(g^{-1}(\gamma)) \frac{\partial g^{-1}(\gamma)}{\partial \gamma}}$$
> $$= \sqrt{I_\gamma(\gamma)}$$

Maximizes the information gained from the data (under appropriate regulatory conditions), i.e. maximizes
$E(KL(f_\theta(.), f_{post}(., x)))$

Empirical Bayes
Let the prior depend on a hyper-parameter: $f_\theta(\theta, \gamma)$
Choose γ s.t. $L(\gamma) = f(x; \gamma) = \int f(x; \vartheta)f_\theta(\vartheta, \gamma)d\vartheta$ is maximal.
Using the data to find the prior contradicts the Bayes approach of incorporating prior knowledge.

Hierarchical Prior

$$x|\theta \sim f(x; \theta); \quad \theta|\gamma \sim f_\theta(\theta, \gamma); \quad \gamma \sim f_\gamma(\gamma)$$

Conjugate Priors
If Prior and Posterior belong to the same family of distributions for a given likelihood function, they are called conjugate.

Examples:

Prior	Likelihood	Posterior
$\pi \sim \text{Be}(\alpha, \beta)$	$\text{Bin}(n, \pi)$	$\text{Be}(\alpha+k, \beta+n-k)$
$\mu \sim \text{N}(\gamma, \tau^2)$	$\text{N}(\mu, \sigma^2)$	$\text{N}(.,.) \overset{n \to \infty}{\longrightarrow} \text{N}(\bar{y}, \frac{\sigma^2}{n})$
$\sigma^2 \sim \text{IG}(\alpha, \beta)$	$\text{N}(\mu, \sigma^2)$	$\text{IG}(\alpha+\frac{n}{2}, \beta+\frac{1}{2}\sum_{i=1}^{n}(y_i-\mu)^2)$
$\lambda \sim \text{Ga}(\alpha, \beta)$	$\text{Po}(\lambda)$	$\text{Ga}(\alpha+n\bar{y}, \beta+n)$

6.2 Numerical Methods for the Posterior

Numerical Integration here: trapezoid approximation
$$\int_\Theta f(y|\vartheta)f_\theta(\vartheta)d\vartheta \approx$$
$$\sum_{k=1}^{K} \frac{f(y; \theta_k)f_\theta(\theta_k) + f(y; \theta_{k-1})f_\theta(\theta_{k-1})}{2}(\theta_k - \theta_{k-1})$$

only normalisation constant unknown, works well for one-dimensional integrals

Laplace Approximation

$$\int_\Theta f(y|\vartheta)f_\theta(\vartheta)d\vartheta \approx f(y; \hat{\theta}_P)f_\theta(\hat{\theta}_P)(2\pi)^{p/2} \left| J_P(\hat{\theta}_P) \right|^{\frac{1}{2}}$$

with the one-dimensional $J_P := -\frac{\partial^2 l_{(n)}(\theta, y)}{\partial \theta^2} - \frac{\partial^2 \log f_\theta(\theta)}{\partial \theta^2}$ Fisher information considering the prior, $\hat{\theta}_P$ posterior mode estimate s.t. $s_{P,\theta}(\hat{\theta}_P) = 0$

Proof:

For n independent samples:

$$f_{post}(\theta|y) = \frac{\prod_{i=1}^{n} f(y_i|\theta) f_\theta(\theta)}{\int \prod_{i=1}^{n} f(y_i|\theta) f_\theta(\theta) d\theta}$$

Denominator: $\int e^{\left\{\sum_{i=1}^{n} \log f(y_i|\theta) + \log f_\theta(\theta)\right\}} d\theta =$

$\int e^{\left\{l(\theta;y) + \log f_\theta(\theta)\right\}} d\theta \overset{TS}{\approx} \int e^{\left(l_P(\hat{\theta}_P) - \frac{1}{2} J_P(\hat{\theta}_P)(\vartheta - \hat{\theta}_P)^2\right)} d\vartheta$

Resembles the normal distribution, therefore the inverse of the normalisation constant can be calculated, which gives the inverse of the Laplace approximation in the univariate case.

Works well for large n and is numerically simple also for big p.

Monte Carlo Approximations

The denominator can be written as $E_\theta(f(y;\theta)) = \int_\Theta f(y|\vartheta) f_\theta(\vartheta) d\vartheta$, which can be estimated by the arithmetic mean for a sample of $\theta_1, ..., \theta_N$, which needs to be drawn from the prior. The following methods to draw from non-standard distributions can be used for that.

- **Inverse CDF**

  $F(X)$ known. Since $F(x) = u$, $F^{-1}(u) = x$, $u \sim U(0,1)$
  1. Draw $u \sim U(0,1)$
  2. Compute $F^{-1}(u)$ to get a value x

Proof:

$$P(x \le y) = P(F^{-1}(u) \le y) = P(u \le F(y)) = F(y)$$

- **Rejection Sampling**

  An umbrella distribution $g(x)$ can be found s.t.
  $\frac{f(x)}{g(x)} \le M \; \forall x$ with $f(x) > 0$ when $g(x) > 0$
  1. Draw candidate $y \sim g(x)$
  2. Acceptance probability α for y: $\alpha = \frac{f(x)}{M g(x)}$
  3. Draw $u \sim U(0,1)$ and accept if $u \le \alpha$, else: step 1

Proof:

$$P\left(Y \le x | U \le \frac{f(Y)}{Mg(Y)}\right) = \frac{P\left(Y \le x, U \le \frac{f(Y)}{Mg(Y)}\right)}{P\left(U \le \frac{f(Y)}{Mg(Y)}\right)}$$

$$= \frac{\int_{-\infty}^{x} \int_0^{\frac{f(y)}{g(x)}} du \, g(y) dy}{\int_{-\infty}^{\infty} \int_0^{\frac{f(y)}{g(x)}} du \, g(y) dy} = \frac{\int_{-\infty}^{x} \frac{f(y)}{g(x)} g(y) dy}{\int_{-\infty}^{\infty} \frac{f(y)}{g(x)} g(y) dy}$$

$$= \frac{\int_{-\infty}^{x} f(y) dy}{\int_{-\infty}^{\infty} f(y) dy} = P(X \le x)$$

- **Importance Sampling**

  Directly estimate $E_\theta(f(y;\theta))$.

  For sampling distribution $g(x)$,

  $$\frac{1}{N} \sum_{i=1}^{n} \frac{f(x)}{g(x)}$$

  is a consistent estimator.

Proof:

$$E_g\left(\frac{1}{N} \sum_{i=1}^{n} \frac{f(x)}{g(x)}\right) = \int \frac{f(x)}{g(x)} g(x) dx = \int f(x) dx = f(x)$$

Markov Chain Monte Carlo   sample from $f_{post}(\theta|X)$

$f(y)$ unknown, however:

$$\frac{f_{post}(\theta|x)}{f_{post}(\tilde{\theta}|x)} = \frac{f(x|\theta) f_\theta(\theta)}{f(y)} \frac{f(y)}{f(x|\tilde{\theta}) f_\theta(\tilde{\theta})} = \frac{f(x|\theta) f_\theta(\theta)}{f(x|\tilde{\theta}) f_\theta(\tilde{\theta})}$$

Metropolis-Hastings: Draw Markov Chain $\theta_1^*, ..., \theta_n^*$:

1. Draw candidate θ^* from proposal distribution $q\left(\theta|\theta_{(t)}^*\right)$

2. Accept $\theta_{(t+1)}^* = \theta^*$ with probability

$$\alpha(\theta_{(t)}|\theta^*) = \min\left\{1, \frac{f_{post}\left(\theta^*|y\right) q\left(\theta_{(t)}^*|\theta^*\right)}{f_{post}\left(\theta_{(t)}^*|y\right) q\left(\theta^*|\theta_{(t)}^*\right)}\right\}$$

else choose $\theta_{(t+1)}^* = \theta_{(t)}^*$

This sequence has a stationary distribution for $n \to \infty$.

Choice of q: trade-off between exploring Θ and reaching a high α.

Burn-in and thinning out give *i.i.d.* samples from $f_{post}(\theta|X)$.

Gibbs Sampling: For high dimensions α is close to zero.

Sample from the marginal distributions seperately:

$$\theta_i^* \sim f_{\theta_i|y, \theta \setminus \theta_i}\left(\theta_i^*|y, \theta_{t^*, i}\right)$$

with $\theta_{t^*, i}$ most recent estimates without θ_i

A Gibbs sampled sequence converges to $f_{post}(\theta|X)$ as stationary.

Can also be used on its own, if marginal densities are known.

Variational Bayes Principles

Approximate $f_{post}(\theta|X)$ by $q_\theta = \min_{q_\theta \in Q} KL(f_{post}(.|X), q_\theta(.))$

Restrict q_θ to independence: $q_\theta(\theta) = \prod_{k=1}^{p} q_k(\theta_k)$

Update each component iteratively. Works well for big p.

7 Sampling

Bootstrap

1. Draw y_i^*: n samples with replacement from y

2. Calculate the statistic of interest $t(y_i^*)$

3. Repeat this B times

4. *Plug-in Principle*: Whenever the distribution function is involved in estimating a statistic, use the empirical bootstrapped version instead.

In a **Parametric Bootstrap** the parameter is first estimated from the data and then Bootstrap samples are drawn from the resulting distribution.

Bootstrap Probability

$$P(Y_i \in Y^*) = 1 - (1 - \frac{1}{n})^n \overset{n \to \infty}{\to} 1 - e^{-1} \approx 0.632$$

Subsampling

- **replacement** m-out-of-n bootstrap

- **non-replacement** subsampling directly from true F

Permutation Test   for two variables

1. Calculate $t(x, y)$, e.g. differences in mean, correlation...

2. Draw samples x^*, y^* of size n from x and y without replacement ("shuffel")

3. Calculate $t(x^*, y^*)$

4. $p - value = \frac{1}{B} \sum_{b=1}^{B} \mathbb{1}_{\{t(x_b^*, y_b^*) \geq t(x,y)\}}$

For a **Boostrap Test** do step 2 with replacement.

Bootstrap in Regression

- **Residual based**: 1. Get Bootstrap sample e_i^* from fitted residuals $\hat{e} = y - X\hat{\beta}$, 2. Calculate new response $y_i^* = x_i\hat{\beta} + e_i^*$, 3. Calculate $\hat{\beta}^*$

- **Model based** 1. Draw a sample from $e_i \sim N(0, \hat{\sigma}^2)$, 2. Calculate new response $y_i^* = x_i\hat{\beta} + e_i^*$, 3. Calculate $\hat{\beta}^*$

- **Pairwise** 1. Draw (y_i^*, x_i^*) from the original sample for $i = 1, .., n$, 2. Calculate $\hat{\beta}^*$

- **Wild** Set $\hat{e}_i^* = V_i^* \hat{e}_i$, with V_i^* from the 2-point distribution $P(V_i^* = \frac{\sqrt{5}+1}{2}) = \frac{\sqrt{5}-1}{2\sqrt{5}}$ and $P(V_i^* = -\frac{\sqrt{5}-1}{2}) = \frac{\sqrt{5}+1}{2\sqrt{5}}$, chosen as $E(V_i^*) = 0$, $Var(V_i^*) = 1$, $E(V_i^{*3}) = 1$

Consistency of a Bootstrap Estimator

$$\lim_{n \to \infty} P_n \left\{ \sup_t |G_n(t, F_n) - G_\infty(t, F)| > \epsilon \right\} = 0 \ \ \forall \epsilon$$

with $F_n(y) = \frac{1}{n} \sum_{i=1}^{n} \mathbb{1}_{\{Y_i \leq y\}}$ empirical distribution function, $G_n(t, F) = P(T_n \leq t)$ exact finite sample distribution, and P_n joint probability of the sample

The bootstrap estimate is inconsistent for the maximum of a sample or if the θ lies on the boundary of Θ.

Mallow's Metric

$$\rho_p(F, G) = \inf_{\mathcal{T}_{XY}} \{E(|X - Y|)^p\}^{\frac{1}{p}}$$

for F, G in the set of distributions where $\inf_{-\infty}^{\infty} |t|^p dF(t) < \infty$; $(X, Y) \sim T \in \mathcal{T}_{XY}$ with $X \sim F$ and $Y \sim G$

Theorem of Beran and Ducharme

$G_n(., F_n)$ is consistent if $\forall \epsilon > 0, F$ the following holds:

1. $\lim_{n \to \infty} P_n(\rho(F_n, F) > \epsilon) = 0$

2. $G_\infty(t, F)$ is a continuous function of t

3. $\forall t$ and sequences $\{H_n\}$ s.t. $\lim_{n \to \infty} \rho(H_n, F) = 0$ holds: $G_n(t, H_n) \to G_\infty(t, F)$

8 Model Selection

AIC (Akaike Information Criterion)

$$AIC = -2\sum_{i=1}^{n}\log f(y_i; \hat{\theta}) + 2p$$

The AIC estimates $2E_Y\{KL(g,f)\} - 2\int \log(g(y))g(y)dy$. The latter component is unknown, so the absolute value of the AIC is not informative. The AIC favours complex models.

For regressions: $AIC = 2n\log(\hat{\sigma}^2) + 2(p+2)$

The AIC as theoretical cross validation

The AIC minimizes $E_Y\left\{E_Y\left[Y - \hat{\mu}\right]^2\right\}$ if we use the MSE instead of the Kullback-Leibler divergence. This can be estimated via cross validation.

Bias Corrected AIC

$$AIC_{corr} = -2\sum_{i=1}^{n}\log f(y_i; \hat{\theta}) + 2p\left(\frac{n}{n-p-1}\right)$$

should be preferred if $\frac{n}{p} < 40$

BIC (Bayesian Information Criterion)

$$BIC = -2\sum_{i=1}^{n}\log f(y_i; \hat{\theta}) + \log(n)p$$

approximately maximizes the posterior probability of a model and selects less complex models as the AIC

DIC (Deviance Information Criterion)   Bayesian AIC

$$DIC = D(y, \hat{\theta}_{postmean}) + 2p_D = \int D(y, \vartheta)f_{post}(\vartheta|y)d\vartheta + p_D$$

with deviance $D(y;\theta) := -2l(\theta)$ the difference in likelihood compared to the full mode and $\Delta D(y;\theta,\hat{\theta}) = 2\left\{l(\hat{\theta}) - l(\theta)\right\} \overset{a}{\sim} \chi_p^2$ the difference in deviance

$$p_D := E(\Delta D(y;\theta,\hat{\theta}_{\substack{post\\mean}}|y) = \int D(y,\vartheta)f_{post}(\vartheta|y)d\vartheta - D(y,\hat{\theta}_{\substack{post\\mean}})$$

The integral can be approximated using MCMC.

Model Averaging   Using probabilities as weights

$$P(M_k|y) := \frac{\exp(-\frac{1}{2}\Delta IC_k)}{\sum_{k'=1}^{K}\exp(-\frac{1}{2}\Delta IC_k')}$$

with $\Delta IC_k = IC_k - min(IC)$

For regression covariates: $P(x_i|y) = \sum_{k=1}^{K}\mathbb{1}_{\{x_i \text{ in } M_k\}}P(M_k|y)$

Inference After Model Selection   neglect is a quiet scandal

$$Var(\hat{\theta}) = E_{model}(Var(\hat{\theta}|model)) + Var_{model}(E(\hat{\theta}|model))$$

$$= \sum_{k=1}^{K}\pi_k Var_k(\hat{\theta}) + \sum_{k=1}^{K}\pi_k(\theta_k - \bar{\theta})^2$$

The last component depends on the true parameter and will be biased if the estimates are used.

Solutions:

- $\widehat{Var}(\hat{\theta}) = \left[\sum_{k=1}^{K}\pi_k\sqrt{\widehat{Var}_k(\hat{\theta}_k) + (\hat{\theta}_k - \hat{\bar{\theta}})^2}\right]^2$

- Use the Variance of the full (saturated) model

- Use bootstrap for confidence intervals

Lasso   least absolute shrinkage and selection operator

$$l_p(\theta, \lambda) = l(\theta) - \lambda\sum_{j=1}^{p}|\theta_j|$$

This penalized log likelihood can be solved with iterative quadratic programming using a Taylor expansion.

Using Bayesian view the penalty corresponds to a prior:

$f_{\theta_j}(\theta_j) \propto \exp(-|\theta_j|)$ (Laplace prior)